

**VIETNAM NATIONAL UNIVERSITY - HO CHI MINH CITY
HO CHI MINH CITY UNIVERSITY OF TECHNOLOGY
FACULTY OF COMPUTER SCIENCE AND ENGINEERING**



MACHINE LEARNING (CO3117)

Assignment Report

Text Classification on AG News

Instructor: Dr. Le Thanh Sach

Members:

Tran Hoang Vy Khang – [2352502]

Le Tan Minh Khoa – [2352563]

Tao Nguyen Quang Khang – [2352499]

Tran Gia Huy – [2252264]

Vo Le Hai Dang – [2352257]

Ho Chi Minh City, Semester I, 2025-2026



Contents

1	Course Context	2
2	Project Objectives	2
3	Dataset and Problem Setting	2
4	Source Code and Reproducibility	2
5	Exploratory Data Analysis (EDA)	2
6	Implemented Methodology	5
6.1	Traditional Pipeline	5
6.2	Modern Embedding Extension	5
7	Results	5
7.1	Best TF-IDF Results (Full Data)	5
7.2	SBERT Benchmark Results	6
7.3	Feature-Family Comparison	6
8	Error Analysis	6
9	Team Contributions	7



1 Course Context

- Course: Machine Learning (CO3117)
- Department: Computer Science and Engineering, HCMUT – VNU-HCM
- Supervisor: Dr. Le Thanh Sach
- Topic: Text Data Classification (Topic 2)

2 Project Objectives

The project aims to build a complete text classification pipeline from exploratory data analysis to final model comparison, following the required workflow:

1. EDA and data understanding.
2. Configurable preprocessing and TF-IDF feature extraction.
3. Classical ML training and evaluation (Logistic Regression, SVM, Naive Bayes).
4. Modern embedding extension with SBERT.
5. Comparison across feature families with report-ready artifacts.

3 Dataset and Problem Setting

- Dataset: ag_news from Hugging Face Datasets.
- Task: 4-class news topic classification.
- Split used in full mode: 120,000 training samples and 7,600 test samples.

4 Source Code and Reproducibility

The complete source code, artifacts, and instructions to run this project are available on our public GitHub repository:

https://github.com/THVKhang/MachineLearning_TextModule

All executions are designed to be reproducible via a single Colab Notebook entry point, without requiring personal cloud storage mounting.

5 Exploratory Data Analysis (EDA)

Extensive Exploratory Data Analysis was performed to understand the characteristics of the AG News dataset.








Figure 1: Overview of Dataset Distributions

Key findings include:

- **Balanced dataset:** The 120,000 training samples are perfectly balanced across the 4 classes (30,000 each). This eliminates the need for class weighting or resampling techniques.
- **Short texts:** The average text length is approximately 35 words. Most semantic content can be captured effectively using TF-IDF with bigrams.
- **Domain-specific vocabulary:** Each class has distinct top keywords (e.g., ‘game’ for Sports, ‘stock’ for Business).
- **Low noise:** The dataset is very clean, with near-zero empty or duplicate texts.
- **Minimal URLs:** Only about 32 URL-heavy texts were found in the training set, having negligible impact on overall quality.



6 Implemented Methodology

6.1 Traditional Pipeline

- Preprocessing options: lowercasing, stopword removal, punctuation removal, number removal.
- TF-IDF configuration search over three candidates: `tfidf_uni_1k`, `tfidf_uni_bi_3k`, and `tfidf_uni_bi_5k`.
- Primary model-selection metric: F1-weighted.
- Trained models: Logistic Regression, SVM, Naive Bayes.

6.2 Modern Embedding Extension

- Embedding model: sentence-transformers/all-MiniLM-L6-v2.
- Benchmark scales: 5k_2k and 20k_2k.
- Embeddings and labels are exported to .npy artifacts.

7 Results

7.1 Best TF-IDF Results (Full Data)

Model	Accuracy	Precision (w)	Recall (w)	F1 (w)
SVM	0.9046	0.9044	0.9046	0.9044
Logistic Regression	0.9038	0.9036	0.9038	0.9036
Naive Bayes	0.8871	0.8866	0.8871	0.8867

Table 1: TF-IDF model comparison.




Figure 3: Confusion Matrices for Traditional Models

7.2 SBERT Benchmark Results

Scale	Best Model	Accuracy	Precision (w)	Recall (w)	F1 (w)
5k_2k	Logistic Regression	0.8740	0.8734	0.8740	0.8736
20k_2k	SVM	0.8955	0.8958	0.8955	0.8954

Table 2: Best embedding model per scale.

7.3 Feature-Family Comparison




Figure 4: Performance Comparison: TF-IDF vs SBERT Embedding

The final comparison shows that the best TF-IDF model (SVM with $F1 = 0.9044$) outperforms the best SBERT embedding model (SVM at 20k scale with $F1 = 0.8954$). TF-IDF is significantly faster to train and extract, while maintaining highly interpretable sparse features. Therefore, TF-IDF + SVM is the primary recommendation.

8 Error Analysis

An error analysis was performed on the best TF-IDF model (SVM) to understand the distribution of misclassifications. Out of 7,600 test samples, the SVM model made 381 errors.

The top 3 most common confusion pairs are:

1. **True: Business** → **Predicted: Sci/Tech** (73 occurrences)
2. **True: Business** → **Predicted: World** (49 occurrences)
3. **True: World** → **Predicted: Sci/Tech** (44 occurrences)

This indicates a semantic overlap between Business/Technology news and Business/World news, which is common in articles discussing international tech corporations or global markets.



9 Team Contributions

Member	ID	Email	Main Contribution	%
Tran Hoang Vy Khang	2352502	2352502@hcmut.edu.vn	Embeddings, workflow, integration	30
Le Tan Minh Khoa	2352563	2352563@hcmut.edu.vn	Text cleaning and TF-IDF	20
Tao Nguyen Quang Khang	2352499	2352499@hcmut.edu.vn	ML training and evaluation	20
Tran Gia Huy	2252264	2252264@hcmut.edu.vn	Pipeline and config architecture	15
Vo Le Hai Dang	2352257	2352257@hcmut.edu.vn	Data loading and EDA analysis	15

Table 3: Detailed contribution table with student emails.