

Machine Learning Course Project Report

Text Classification on AG News

Tran Hoang Vy Khang (2352502)
Le Tan Minh Khoa (2352563)
Tao Nguyen Quang Khang (2352499)
Tran Gia Huy (2252264)
Vo Le Hai Dang (2352257)

Semester I, 2025–2026

1 Course Context

- Course: Machine Learning (CO3117)
- Department: Computer Science and Engineering, HCMUT – VNU-HCM
- Supervisor: Dr. Le Thanh Sach
- Topic: Text Data Classification (Topic 2)

2 Project Objectives

The project aims to build a complete text classification pipeline from exploratory data analysis to final model comparison, following the required workflow:

1. EDA and data understanding.
2. Configurable preprocessing and TF-IDF feature extraction.
3. Classical ML training and evaluation (Logistic Regression, SVM, Naive Bayes).
4. Modern embedding extension with SBERT.
5. Comparison across feature families with report-ready artifacts.

3 Dataset and Problem Setting

- Dataset: `ag_news` from Hugging Face Datasets.
- Task: 4-class news topic classification.
- Split used in full mode: 120,000 training samples and 7,600 test samples.

4 Source Code and Reproducibility

The complete source code, artifacts, and instructions to run this project are available on our public GitHub repository:

https://github.com/THVKhang/MachineLearning_TextModule

All executions are designed to be reproducible via a single Colab Notebook entry point, without requiring personal cloud storage mounting.

5 Exploratory Data Analysis (EDA)

Extensive Exploratory Data Analysis was performed to understand the characteristics of the AG News dataset.




Figure 1: Overview of Dataset Distributions

Key findings include:

- **Balanced dataset:** The 120,000 training samples are perfectly balanced across the 4 classes (30,000 each). This eliminates the need for class weighting or resampling techniques.
- **Short texts:** The average text length is approximately 35 words. Most semantic content can be captured effectively using TF-IDF with bigrams.
- **Domain-specific vocabulary:** Each class has distinct top keywords (e.g., ‘game” for Sports, ‘stock” for Business).
- **Low noise:** The dataset is very clean, with near-zero empty or duplicate texts.
- **Minimal URLs:** Only about 32 URL-heavy texts were found in the training set, having negligible impact on overall quality.




Figure 2: Word Frequency Analysis across AG News

6 Implemented Methodology

6.1 Traditional Pipeline

- Preprocessing options: lowercasing, stopwords removal, punctuation removal, number removal.
- TF-IDF configuration search over three candidates: `tfidf_uni_1k`, `tfidf_uni_bi_3k`, and `tfidf_uni_bi_5k`.
- Primary model-selection metric: F1-weighted.
- Trained models: Logistic Regression, SVM, Naive Bayes.

6.2 Modern Embedding Extension

- Embedding model: `sentence-transformers/all-MiniLM-L6-v2`.
- Benchmark scales: `5k_2k` and `20k_2k`.
- Embeddings and labels are exported to `.npy` artifacts.

7 Results

7.1 Best TF-IDF Results (Full Data)

Model	Accuracy	Precision (w)	Recall (w)	F1 (w)
SVM	0.9046	0.9044	0.9046	0.9044
Logistic Regression	0.9038	0.9036	0.9038	0.9036
Naive Bayes	0.8871	0.8866	0.8871	0.8867

Table 1: TF-IDF model comparison.




Figure 3: Confusion Matrices for Traditional Models

7.2 SBERT Benchmark Results

Scale	Best Model	Accuracy	Precision (w)	Recall (w)	F1 (w)
5k_2k	Logistic Regression	0.8740	0.8734	0.8740	0.8736
20k_2k	SVM	0.8955	0.8958	0.8955	0.8954

Table 2: Best embedding model per scale.

7.3 Feature-Family Comparison




Figure 4: Performance Comparison: TF-IDF vs SBERT Embedding

The final comparison shows that the best TF-IDF model (SVM with $F1 = 0.9044$) outperforms the best SBERT embedding model (SVM at 20k scale with $F1 = 0.8954$). TF-IDF is significantly faster to train and extract, while maintaining highly interpretable sparse features. Therefore, TF-IDF + SVM is the primary recommendation.

8 Error Analysis

An error analysis was performed on the best TF-IDF model (SVM) to understand the distribution of misclassifications. Out of 7,600 test samples, the SVM model made 381 errors.

The top 3 most common confusion pairs are:

1. **True: Business** → **Predicted: Sci/Tech** (73 occurrences)
2. **True: Business** → **Predicted: World** (49 occurrences)
3. **True: World** → **Predicted: Sci/Tech** (44 occurrences)

This indicates a semantic overlap between Business/Technology news and Business/World news, which is common in articles discussing international tech corporations or global markets.

9 Team Contributions

Member	ID	Email	Main Contribution	%
Tran Hoang Vy Khang	2352502	2352502@hcmut.edu.vn	Embeddings, workflow, integration	30
Le Tan Minh Khoa	2352563	2352563@hcmut.edu.vn	Text cleaning and TF-IDF	20
Tao Nguyen Quang Khang	2352499	2352499@hcmut.edu.vn	ML training and evaluation	20
Tran Gia Huy	2252264	2252264@hcmut.edu.vn	Pipeline and config architecture	15
Vo Le Hai Dang	2352257	2352257@hcmut.edu.vn	Data loading and EDA analysis	15

Table 3: Detailed contribution table with student emails.